

MPI File Transfer Program Report

Hoang Minh Quan
ID 23BI14371

December 5, 2025

Contents

1	Introduction	2
2	Program Objectives	2
3	System Illustration	2
4	Source Code	2
5	How to Compile	6
6	How to Run the Program	6
7	Conclusion	6

1 Introduction

This report presents an MPI-based file transfer program implemented in C++. The goal of the program is to allow one MPI process (the ROOT process) to read a file and distribute its contents to all other MPI processes using the `MPI_Bcast` communication primitive.

Each non-root process receives the file and writes it to a new local file named according to its MPI rank.

2 Program Objectives

The application fulfills the following objectives:

- The ROOT process reads the name and contents of a file.
- The file name and its length are broadcast to all MPI processes.
- The file data buffer is broadcast to all MPI processes.
- Each receiver process writes the received data into a new file named:

`received_rankX_filename`

The program requires at least 2 MPI processes:

- One sender (rank 0)
- One or more receivers (ranks 1, 2, ...)

3 System Illustration

Figure 1 shows a high-level view of the MPI file transfer system. Rank 0 acts as the sender: it reads the input file and broadcasts its contents to all other ranks, which act as receivers.

4 Source Code

The full source code of the MPI program is shown below.




Figure 1: High-level illustration of the MPI file transfer system.

```

1 #include <mpi.h>
2 #include <iostream>
3 #include <fstream>
4 #include <vector>
5 #include <string>
6
7 int main(int argc, char** argv) {
8     MPI_Init(&argc, &argv);
9
10    int rank, size;
11    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
12    MPI_Comm_size(MPI_COMM_WORLD, &size);
13
14    const int ROOT = 0;
15
16    if (size < 2) {
17        if (rank == ROOT) {
18            std::cerr << "Need at least 2 MPI processes (1
19                         sender + 1 receiver)." << std::endl;
20        }
21        MPI_Finalize();
22        return 1;
23    }
24
25    std::string filename;
26    std::vector<char> nameBuf;
27    int nameLen = 0;
28
29    if (rank == ROOT) {
30        if (argc < 2) {

```

```

30         std::cerr << "Usage: mpirun -np <n> ./"
31             mpi_file_transfer <file_to_send>\n";
32         MPI_Abort(MPI_COMM_WORLD, 1);
33     }
34     filename = argv[1];
35
36     nameBuf.assign(filename.begin(), filename.end());
37     nameBuf.push_back('\0');
38     nameLen = static_cast<int>(nameBuf.size());
39 }
40
41 MPI_Bcast(&nameLen, 1, MPI_INT, ROOT, MPI_COMM_WORLD);
42
43 if (rank != ROOT) {
44     nameBuf.resize(nameLen);
45 }
46
47 MPI_Bcast(nameBuf.data(), nameLen, MPI_CHAR, ROOT,
48 MPI_COMM_WORLD);
49
50 if (rank != ROOT) {
51     filename = std::string(nameBuf.data());
52 }
53
54 int fileSize = 0;
55 std::vector<char> buffer;
56
57 if (rank == ROOT) {
58     std::ifstream in(filename, std::ios::binary);
59     if (!in.is_open()) {
60         std::cerr << "Rank " << rank << ": cannot open
61             file '" << filename << "' for reading.\n";
62         MPI_Abort(MPI_COMM_WORLD, 1);
63     }
64
65     in.seekg(0, std::ios::end);
66     fileSize = static_cast<int>(in.tellg());
67     in.seekg(0, std::ios::beg);
68
69     buffer.resize(fileSize);
70     if (fileSize > 0) {
71         in.read(buffer.data(), fileSize);
72     }

```

```

70     in.close();
71
72     std::cout << "Rank " << rank << ": read " <<
73         fileSize
74             << " bytes from '" << filename << "'.\n";
75
76     MPI_Bcast (&fileSize, 1, MPI_INT, ROOT, MPI_COMM_WORLD);
77
78     if (fileSize == 0) {
79         if (rank != ROOT) {
80             std::cerr << "Rank " << rank
81                 << ": received file size 0, nothing
82                     to write.\n";
83         }
84         MPI_Finalize();
85         return 0;
86     }
87
88     if (rank != ROOT) {
89         buffer.resize(fileSize);
90     }
91
92     MPI_Bcast (buffer.data(), fileSize, MPI_CHAR, ROOT,
93                 MPI_COMM_WORLD);
94
95     if (rank != ROOT) {
96         std::string outName = "received_rank"
97                         + std::to_string(rank)
98                         + "_" + filename;
99
100        std::ofstream out(outName, std::ios::binary);
101        if (!out.is_open()) {
102            std::cerr << "Rank " << rank
103                << ": cannot open '" << outName << "'"
104                    for writing.\n";
105        MPI_Abort (MPI_COMM_WORLD, 1);
106    }
107
108    out.write (buffer.data(), fileSize);
109    out.close();
110
111    std::cout << "Rank " << rank << ": wrote " <<

```

```

109         fileSize
110             << " bytes to '" << outName << "'.\n";
111     }
112     MPI_Finalize();
113     return 0;
114 }
```

Listing 1: MPI File Transfer Source Code

5 How to Compile

To compile the program on a system with MPI installed:

```
mpic++ mpi_file_transfer.cpp -o mpi_file_transfer
```

6 How to Run the Program

Run the program with at least 2 processes:

```
mpirun -np 4 ./mpi_file_transfer example.bin
```

Expected behavior:

- Rank 0 reads `example.bin`.
- All ranks receive the file through MPI broadcast.
- Ranks 1, 2, 3 write:

```

received_rank1_example.bin
received_rank2_example.bin
received_rank3_example.bin
```

7 Conclusion

This MPI application demonstrates how to efficiently distribute file data from one process to all others using collective communication.

The program:

- Uses MPI_Bcast for distributing metadata and file contents.
- Avoids sending the file multiple times by using a single collective operation.
- Ensures each process receives an identical copy of the input file.

This approach is well-suited for small to medium file sizes and showcases a clean example of using MPI for data distribution.